

## Keras -- MLPs on MNIST

In [0]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

In [0]:

```
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, colors=['b']):
    plt.plot(x, vy, color = 'b', label='Validation Loss')
    plt.plot(x, ty, color = 'r', label='Train Loss')
    plt.xlabel('epoch')
    plt.ylabel('Categorical Crossentropy Loss')
    plt.legend()
    plt.grid()
    plt.show();
```

In [0]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 1s 0us/step

In [0]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
Number of training examples : 10000 and each image is of shape (28, 28)

In [0]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [0]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)  
Number of training examples : 10000 and each image is of shape (784)

In [0]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

In [0]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

In [0]:

```
# example data point after normlizing
print(X_train[0])
```

```
[0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0  0  0  0  0  0]
```

[illegible]

[illegible]

In [0]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

## Softmax classifier

In [0]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, BatchNormalization, Activation
```

In [0]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

## 1. two hidden layers

In [0]:

```
# start building a model
model = Sequential()
```

```
# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input
```

```
# output_dim represent the number of nodes need in that layer
# here we have 10 nodes
```

```
model.add(Dense(units=256, input_shape=(input_dim,), activation='relu'))
model.add(Dropout(0.25))
```

```
model.add(Dense(128))
model.add(BatchNormalization())
model.add(Activation('relu'))
```

```
model.add(Dense(output_dim, activation='softmax'))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:4432: The name tf.random\_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:148: The name tf.placeholder\_with\_default is deprecated. Please use tf.compat.v1.placeholder\_with\_default instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3733: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

In [0]:

```
model.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 256)	200960
dropout_1 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 128)	32896
batch_normalization_1 (Batch Normalization)	(None, 128)	512
activation_1 (Activation)	(None, 128)	0
dense_5 (Dense)	(None, 10)	1290
Total params: 235,658		
Trainable params: 235,402		
Non-trainable params: 256		

In [0]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method
```

```
# It receives three arguments:
```

```
# An optimizer. This could be the string identifier of an existing optimizer, https://keras.io/optimizers/
```

```
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
```

```
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/
```

```

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all
# zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted our labels into vectors

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per
# _epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values
# (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_dat
a=(X_test, Y_test))

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow\_core/python/ops/math\_grad.py:1424: where (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:1033: The name tf.assign\_add is deprecated. Please use tf.compat.v1.assign\_add instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3005: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:190: The name tf.get\_default\_session is deprecated. Please use tf.compat.v1.get\_default\_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:207: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:216: The name tf.is\_variable\_initialized is deprecated. Please use tf.compat.v1.is\_variable\_initialized instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:223: The name tf.variables\_initializer is deprecated. Please use tf.compat.v1.variables\_initializer instead.

60000/60000 [=====] - 13s 214us/step - loss: 0.2877 - acc: 0.9166 - val\_loss: 0.1143 - val\_acc: 0.9649

Epoch 2/20

60000/60000 [=====] - 4s 65us/step - loss: 0.1233 - acc: 0.9627 - val\_loss: 0.0876 - val\_acc: 0.9733

Epoch 3/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0895 - acc: 0.9723 - val\_loss: 0.0763 - val\_acc: 0.9774

Epoch 4/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0876 - acc: 0.9774 - val\_loss: 0.0763 - val\_acc: 0.9774

```

60000/60000 [=====] - 4s 65us/step - loss: 0.0726 - acc: 0.9771 - val_loss: 0.0655 - val_acc: 0.9796
Epoch 5/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0631 - acc: 0.9797 - val_loss: 0.0640 - val_acc: 0.9802
Epoch 6/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0536 - acc: 0.9831 - val_loss: 0.0668 - val_acc: 0.9797
Epoch 7/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0470 - acc: 0.9845 - val_loss: 0.0676 - val_acc: 0.9798
Epoch 8/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0454 - acc: 0.9854 - val_loss: 0.0634 - val_acc: 0.9798
Epoch 9/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0395 - acc: 0.9870 - val_loss: 0.0683 - val_acc: 0.9790
Epoch 10/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0393 - acc: 0.9868 - val_loss: 0.0569 - val_acc: 0.9826
Epoch 11/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0311 - acc: 0.9894 - val_loss: 0.0608 - val_acc: 0.9811
Epoch 12/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0309 - acc: 0.9894 - val_loss: 0.0626 - val_acc: 0.9826
Epoch 13/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0283 - acc: 0.9902 - val_loss: 0.0618 - val_acc: 0.9818
Epoch 14/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0274 - acc: 0.9903 - val_loss: 0.0610 - val_acc: 0.9811
Epoch 15/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0273 - acc: 0.9907 - val_loss: 0.0597 - val_acc: 0.9826
Epoch 16/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0250 - acc: 0.9916 - val_loss: 0.0637 - val_acc: 0.9823
Epoch 17/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0237 - acc: 0.9915 - val_loss: 0.0576 - val_acc: 0.9833
Epoch 18/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0219 - acc: 0.9925 - val_loss: 0.0646 - val_acc: 0.9830
Epoch 19/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0206 - acc: 0.9929 - val_loss: 0.0611 - val_acc: 0.9822
Epoch 20/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0229 - acc: 0.9921 - val_loss: 0.0633 - val_acc: 0.9821

```

In [0]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']


```



Test score: 0.063285964876238  
Test accuracy: 0.9821

In [0]:

```
plt_dynamic(x, vy, ty)
```



## 2. three hidden layer

In [0]:

```
# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(units=256, input_shape=(input_dim,), activation='relu'))
model.add(Dropout(0.25))

model.add(Dense(units=256, activation='relu'))
model.add(Dropout(0.25))

model.add(Dense(128))
model.add(BatchNormalization())
model.add(Activation('relu'))

model.add(Dense(output_dim, activation='softmax'))
```

In [0]:

```
model.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 256)	200960
dropout_2 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 256)	65792
dropout_3 (Dropout)	(None, 256)	0
dense_8 (Dense)	(None, 128)	32896

batch_normalization_2 (Batch Normalization)	(None, 128)	512
activation_2 (Activation)	(None, 128)	0
dense_9 (Dense)	(None, 10)	1290
=====		
Total params: 301,450		
Trainable params: 301,194		
Non-trainable params: 256		
=====		

In [0]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer , https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values
(if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 5s 80us/step - loss: 0.3157 - acc: 0.9054 - val_loss: 0.1099 - val_acc: 0.9668
Epoch 2/20
60000/60000 [=====] - 4s 73us/step - loss: 0.1349 - acc: 0.9589 - val_loss: 0.0921 - val_acc: 0.9710
Epoch 3/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1037 - acc: 0.9682 - val_loss: 0.0806 - val_acc: 0.9756
Epoch 4/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0863 - acc: 0.9732 - val_loss: 0.0690 - val_acc: 0.9798
Epoch 5/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0741 - acc: 0.9765 - val_loss: 0.0660 - val_acc: 0.9807
Epoch 6/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0650 - acc: 0.9792 - val_loss: 0.0669 - val_acc: 0.9801
Epoch 7/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0567 - acc: 0.9817 - val_loss: 0.0669 - val_acc: 0.9801
```

```

60000/60000 [-----] - 4s 71us/step - loss: 0.0367 - acc: 0.9817 - val_loss: 0.
0671 - val_acc: 0.9807
Epoch 8/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0533 - acc: 0.9826 - val_loss: 0.
0682 - val_acc: 0.9808
Epoch 9/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0489 - acc: 0.9838 - val_loss: 0.
0673 - val_acc: 0.9811
Epoch 10/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0440 - acc: 0.9864 - val_loss: 0.
0680 - val_acc: 0.9807
Epoch 11/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0422 - acc: 0.9864 - val_loss: 0.
0669 - val_acc: 0.9813
Epoch 12/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0398 - acc: 0.9870 - val_loss: 0.
0680 - val_acc: 0.9820
Epoch 13/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0363 - acc: 0.9883 - val_loss: 0.
0602 - val_acc: 0.9836
Epoch 14/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0322 - acc: 0.9890 - val_loss: 0.
0620 - val_acc: 0.9831
Epoch 15/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0344 - acc: 0.9882 - val_loss: 0.
0583 - val_acc: 0.9817
Epoch 16/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0303 - acc: 0.9899 - val_loss: 0.
0648 - val_acc: 0.9828
Epoch 17/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0299 - acc: 0.9900 - val_loss: 0.
0686 - val_acc: 0.9819
Epoch 18/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0286 - acc: 0.9904 - val_loss: 0.
0619 - val_acc: 0.9829
Epoch 19/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0266 - acc: 0.9913 - val_loss: 0.
0627 - val_acc: 0.9835
Epoch 20/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0270 - acc: 0.9912 - val_loss: 0.
0567 - val_acc: 0.9850

```

In [0]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']


```

Test score: 0.05666124046119221

Test accuracy: 0.985

In [0]:

```
plt_dynamic(x, vy, ty)
```



### 3. five hidden layer

In [0]:

```
# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(units=256, input_shape=(input_dim,), activation='relu'))
model.add(Dropout(0.25))

model.add(Dense(units=256, activation='relu'))
model.add(Dropout(0.25))

model.add(Dense(units=512, activation='relu'))
model.add(Dropout(0.5))

model.add(Dense(units=256, activation='relu'))
model.add(Dropout(0.25))

model.add(Dense(128))
model.add(BatchNormalization())
model.add(Activation('relu'))

model.add(Dense(output_dim, activation='softmax'))
```

In [0]:

```
model.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 256)	200960
dropout_8 (Dropout)	(None, 256)	0
dense_17 (Dense)	(None, 256)	65792
dropout_9 (Dropout)	(None, 256)	0
dense_18 (Dense)	(None, 512)	131584
dropout_10 (Dropout)	(None, 512)	0
dense_19 (Dense)	(None, 256)	131008

dense_19 (Dense)	(None, 256)	131328
dropout_11 (Dropout)	(None, 256)	0
dense_20 (Dense)	(None, 128)	32896
batch_normalization_4 (Batch Normalization)	(None, 128)	512
activation_4 (Activation)	(None, 128)	0
dense_21 (Dense)	(None, 10)	1290
Total params: 564,362		
Trainable params: 564,106		
Non-trainable params: 256		

In [0]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer , https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all
# zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted our labels into vectors

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per
# _epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values
# (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 6s 103us/step - loss: 0.4301 - acc: 0.8676 - val_loss: 0.1490 - val_acc: 0.9553
Epoch 2/20
60000/60000 [=====] - 5s 85us/step - loss: 0.1758 - acc: 0.9497 - val_loss: 0.1054 - val_acc: 0.9684
Epoch 3/20
60000/60000 [=====] - 5s 86us/step - loss: 0.1336 - acc: 0.9618 - val_loss: 0.0877 - val_acc: 0.9751
Epoch 4/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1079 - acc: 0.9682 - val_loss: 0.0876 - val_acc: 0.9758
Epoch 5/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0968 - acc: 0.9719 - val_loss: 0.0837 - val_acc: 0.9764
```

```

Epoch 6/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0852 - acc: 0.9759 - val_loss: 0.0800 - val_acc: 0.9777
Epoch 7/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0804 - acc: 0.9762 - val_loss: 0.0677 - val_acc: 0.9805
Epoch 8/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0717 - acc: 0.9795 - val_loss: 0.0703 - val_acc: 0.9804
Epoch 9/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0652 - acc: 0.9808 - val_loss: 0.0771 - val_acc: 0.9790
Epoch 10/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0639 - acc: 0.9817 - val_loss: 0.0739 - val_acc: 0.9799
Epoch 11/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0552 - acc: 0.9834 - val_loss: 0.0820 - val_acc: 0.9794
Epoch 12/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0530 - acc: 0.9847 - val_loss: 0.0704 - val_acc: 0.9801
Epoch 13/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0513 - acc: 0.9850 - val_loss: 0.0687 - val_acc: 0.9817
Epoch 14/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0450 - acc: 0.9866 - val_loss: 0.0743 - val_acc: 0.9808
Epoch 15/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0464 - acc: 0.9861 - val_loss: 0.0696 - val_acc: 0.9816
Epoch 16/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0464 - acc: 0.9862 - val_loss: 0.0671 - val_acc: 0.9831
Epoch 17/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0391 - acc: 0.9888 - val_loss: 0.0779 - val_acc: 0.9796
Epoch 18/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0399 - acc: 0.9878 - val_loss: 0.0690 - val_acc: 0.9835
Epoch 19/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0390 - acc: 0.9882 - val_loss: 0.0624 - val_acc: 0.9841
Epoch 20/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0326 - acc: 0.9904 - val_loss: 0.0728 - val_acc: 0.9827

```

In [0]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']

```

Test score: 0.07284883902525763  
Test accuracy: 0.9827

In [0]:

```
plt_dynamic(x, vy, ty)
```



In [0]:

```
from prettytable import PrettyTable
```

In [0]:

```
x = PrettyTable()
x.field_names = ["#layers", "Train Loss/Acc", "Test Loss/Acc"]
```

In [0]:

```
x.add_row(['2 layers (256 (p=25%), 128 (BN))', '2.3% / 99.21%', '6.3% / 98.21%'])
x.add_row(['3 layers (256 (p=25%), 256 (p=25%), 128 (BN))', '2.7% / 99.12%', '5.7% / 98.5%'])
x.add_row(['5 layers (256 (p=25%), 256 (p=25%), 512 (p=50%), 256 (p=25%), 128 (BN))', '3.3% / 99.04%', '7.3% / 98.27%'])
```

In [99]:

```
print(x)
```

#layers	Train Loss/Acc	Test Loss/Acc
2 layers (256 (p=25%), 128 (BN))	2.3% / 99.21%	6.3% / 98.21%
3 layers (256 (p=25%), 256 (p=25%), 128 (BN))	2.7% / 99.12%	5.7% / 98.5%
5 layers (256 (p=25%), 256 (p=25%), 512 (p=50%), 256 (p=25%), 128 (BN))	3.3% / 99.04%	7.3% / 98.27%

In [0]: